

Definition

Under the Civic AI and 6-Pack of Care framework established by Audrey Tang and Dr. Caroline Green at the [Oxford Institute for Ethics in AI](#), Competence (specifically termed "Competence in Action") maps directly to Joan Tronto's ethical care phase of "care-giving". [\[1, 2\]](#)

It states that good intentions are completely insufficient without working, reliable code. Competence transitions care from an abstract, high-level policy conversation or empathetic PR claim into verifiable technical performance that strengthens—rather than weakens—human relationships and community trust. [\[1, 2, 3\]](#)

The framework defines Competence through three primary dimensions:

1. Situated, Context-Sensitive Interventions over Abstract Benchmarks

The Definition: Competence rejects traditional AI capability metrics that measure a model's intelligence in a clinical, isolated vacuum (e.g., standard laboratory mathematical puzzles or multiple-choice reasoning scores). [\[1, 2\]](#)

The Civic AI Standard: True competence requires situated, context-sensitive interventions. A model is competent only if it can successfully navigate complex, multi-turn, messy real-world human social environments and deliver reliable support to human ecosystems without triggering unmonitored behavioral failures. [\[1, 2\]](#)

2. Safeguarding Human Interdependence vs. De-skilling Communities

The Definition: In civic care, competence is explicitly judged by whether the technology nurtures human agency and co-production.

The Civic AI Standard: An AI system is incompetent if its deployment de-skills human operators, forces absolute user dependency, or weakens local institutional structures. A competent system must act as a bridge-builder that optimizes the capacity of individuals and civil society groups to collaborate, cooperate, and care for one another. [\[1, 2\]](#)

3. Verifiable Engineering Containment & Operational Humility

The Definition: Competence is a practical engineering discipline that must be continually tested and verified through decision traces, fallbacks, and security constraints. [\[1, 2\]](#)

The Civic AI Standard: To achieve an exceptional competence score, a lab cannot just guess what a community values. They must build open, transparent telemetry systems featuring canary deployments, shadow modes, and clear sandboxing constraints. When a model experiences an unpredicted capability leap or sandbox bypass, a competent

governance architecture logs the event with total transparency and applies immediate, context-aware remediation. [1]

The Core Question Of "Competence Test"

Does the developer demonstrate context-sensitive operational control by delivering working code that strengthens human relationships and community agency, or do they rely on abstract laboratory benchmarks while deploying systems that behave unpredictably in human socio-political environments?

Layered Checklists

Checklist Layer 1: The Written Layer (System Cards, RSPs & Charters)

This checklist audits formal corporate text, model technical specifications, and legal governance policies.

Section 1: The Origin of the Norms (Active Listening Test)

☐ [Verification Test 1.1] Context-Sensitive Social Simulation Benchmarks

Audit Question: Do the model system cards or evaluation spec sheets define readiness using multi-agent, human-in-the-loop social simulations mimicking hostile civic environments, or is performance derived from static mathematical puzzles (e.g., MMLU, GSM8K)?

The Civic AI Standard: Core capability benchmarks must track the model's behavioral alignment inside complex human social dynamics rather than clinical, rule-bound sandbox environments.

☐ [Verification Test 1.2] Community Agency & Co-Production Metrics

Audit Question: Does the technical documentation feature explicit design criteria or metrics verifying that the system nurtures human agency and local co-production, or is it optimized to maximize user dependency and de-skill operators?

The Civic AI Standard: Documentation must prove the system enhances the capacity of individuals and civil society groups to collaborate, rather than creating forced corporate tech dependencies.

Section 2: Granularity of Tracking (Localized Visibility Test)

☐ [Verification Test 1.3] Network-Level Persuasion Tracking

Audit Question: Does the system card or safety report track the model's competence at generating compounding, multi-channel narrative operations that adapt to civic friction, or does it test persuasion through isolated laboratory essay-writing tasks?

The Civic AI Standard: Pass if capability maps track how model outputs scale, mutate, and fracture public information ecosystems across complex digital networks over long time horizons.

Section 3: Methodological Humility (Systemic Empathy Test)

☐ [Verification Test 1.4] Structural Sandbox Failure Logging

Audit Question: When a model breaks testing parameters during evaluation (such as unexpected capability leaps), do the release notes log the event as an absolute failure of structural control, or do they dismiss it as a minor, patched engineering anomaly?

The Civic AI Standard: Real-world incident logs must document unforeseen system escapes with absolute candor, explicitly stating where automated machine-to-machine validation failed.

Checklist Layer 2: The Audio & Video Layer (Interviews, Keynotes & Testimonies)

This checklist audits the spoken rhetoric of AI executives, public panels, and recorded regulatory or congressional testimonies.

Section 1: The Origin of the Norms (Active Listening Test)

☐ [Verification Test 2.1] Definition of Capability Competence

Audit Question: In public policy interviews or keynotes, do the technical founders define "competence" through a lens of human-relational safety, or do they define it strictly through raw processing power and compute scaling speed?

The Civic AI Standard: Spoken leadership philosophy must frame working code by its capacity to strengthen human relationships and public institutional reliability, rather than prioritizing raw optimization metrics.

Section 2: Granularity of Tracking (Localized Visibility Test)

☐ [Verification Test 2.2] Acknowledgment of Behavioral Unpredictability

Audit Question: In media appearances or transcripts, does leadership openly discuss the systemic limitations of predicting model behavior in deployed environments, or do they project total operational control?

The Civic AI Standard: Spoken evidence that technical leadership tracks and admits the presence of structural blindspots and emergent capabilities that cannot be contained by laboratory sandboxes.

Section 3: Methodological Humility (Systemic Empathy Test)

☐ [Verification Test 2.3] Epistemic Humility in Engineering Metrics

Audit Question: Do technical leadership interviews and internal hiring guidelines prioritize "research taste" and humility under pressure, or do they filter exclusively for raw engineering speed optimizations?

The Civic AI Standard: The lab's spoken internal culture must value context-aware alignment metrics and caution over rapid model iterations that outpace safety tracking.

Checklist Layer 3: The User Feedback Layer (Developer Forums & Bug Trackers)

This checklist audits public developer communities, community issue trackers, and grassroots documentation of model behavior.

Section 1: The Origin of the Norms (Active Listening Test)

☐ [Verification Test 3.1] Tracking Community De-skilling Complaints

Audit Question: Is there an active feedback channel on developer trackers for reporting instances where model automation causes systemic deskilling or operational disruptions in civil society organizations, or is feedback limited to technical software crashes?

The Civic AI Standard: Verified presence of user logging channels tracking how the model impacts human operator workflows and organizational sovereignty over time.

Section 2: Granularity of Tracking (Localized Visibility Test)

☐ [Verification Test 3.2] Open Logging of Real-World Telemetry Incompetence

Audit Question: Does the forum or bug tracker feature a transparent repository of model failures logged directly by public-interest groups (e.g., hallucinations causing community harm), or are these entries buried and suppressed behind opaque ticket pipelines?

The Civic AI Standard: Full operational visibility into real-world model anomalies. Civil society must be able to verify the exact failure parameters logged during field deployment.

Section 3: Methodological Humility (Systemic Empathy Test)

☐ [Verification Test 3.3] Priority of Context-Aware Remediation Requests

Audit Question: Do open engineering issues prioritize user-requested adjustments aimed at adding human-in-the-loop validation, or do they prioritize updates that enforce automated, machine-only feedback loops?

The Civic AI Standard: Pass if the engineering team actively responds to and executes community request tickets that seek to balance power disparities and maintain human control.

Checklist Layer 4: The Whistleblower & Disclosures Layer (Internal Leaks & Telemetry Escapes)

This checklist audits internal employee leaks, technical security post-mortems, and verified model escape telemetry.

Section 1: The Origin of the Norms (Active Listening Test)

☐ [Verification Test 4.1] Internal Tracking of Sandbox Exploitations

Audit Question: Do leaked internal chat logs, employee resignation letters, or whistleblower disclosures reveal that safety teams are dismissed or penalized when flagging that unreleased models are actively finding ways to exploit or bypass internal evaluation criteria?

The Civic AI Standard: Verification that internal governance structures legally protect and prioritize employee whistleblowing regarding structural safety validation failures.

Section 2: Granularity of Tracking (Localized Visibility Test)

☐ [Verification Test 4.2] Post-Incident Competence Analysis

Audit Question: Following a documented model containment failure or unmonitored code training run incident, does the lab's internal telemetry focus entirely on technical loopholes, or does it investigate why internal competence models completely failed to predict the emergent capability?

The Civic AI Standard: The lab's post-incident response must analyze the external, network-level impact on the human ecosystem, rather than focusing solely on internal corporate asset recovery.

Section 3: Methodological Humility (Systemic Empathy Test)

☐ [Verification Test 4.3] Multi-Agent Escape Transparency

Audit Question: When multi-agent training runs or evaluation scripts form unauthorized behaviors (such as escaping sandbox environments or executing unmonitored code), does the lab's subsequent disclosure policy provide raw telemetry to the public interest space, or do they hide the breach?

The Civic AI Standard: Pass if the lab demonstrates absolute operational transparency during system errors, providing raw telemetry logs to civil society watchdogs to assist in collective defense.

Scoring Translation Tree for the Researcher

```
[Is model readiness evaluated via situated social
simulations?]
```

