

Definition

In the Civic AI framework developed by Audrey Tang and Dr. Caroline Green at the [Oxford Institute for Ethics in AI](#), Responsibility (specifically termed "Responsibility in Engagement") is the second core pillar of the [6-Pack of Care](#). [1, 2]

It adapts political philosopher Joan Tronto's ethical phase of "caring for"—taking on the active burden to address a recognized need—and translates it into explicit engineering and governance constraints. [1, 2]

The framework defines Responsibility through three main operational dimensions:

1. Sustained Engagement vs. Rigid Top-Down Control

The Definition: Responsibility is not defined as a sterile compliance checklist or an unyielding corporate veto system. Instead, it requires the AI system and its developers to maintain an active, flexible engagement with dynamic human relational contexts. []

The Core Concept: Rather than treating alignment as a static, one-time moral upload (like hardcoding a permanent corporate baseline), Responsibility treats accountability as an ongoing relationship that must continuously adjust to community feedback. [1]

2. Relational Commitments over Liability Offloading

The Definition: Responsibility demands clear, public-facing accountability structures that define exactly who is accountable, with what authority, and what material consequences happen if the system fails. [1, 2]

The Core Concept: It strictly rejects the corporate practice of downward liability shifting, where frontier labs use boilerplate End User License Agreements (EULAs) to place 100% of the legal and moral risks on end-users or civil society developers. Responsibility requires developers to form "credible commitments" and assume shared structural stewardship alongside the community. [1]

3. Binding Accountability via Alignment Assemblies

The Definition: In practice, Responsibility requires corporate safety assertions to be contractually and operationally bound to public consensus.

The Core Concept: True responsibility cannot exist through voluntary corporate promises or internal board configurations alone. It is established when a lab ties its deployment milestones to public pledges and democratic alignment assemblies (such as [Taiwan's](#)

[deliberative civic frameworks](#)), giving external public-interest stakeholders actual governing power over the AI system. [, [2](#), [3](#)]

The Core Question Of "Responsibility Test"

Does the model developer assume shared legal, financial, and structural liability through community-facing commitments, or do they insulate the corporation by shifting the operational risks of model failure onto civil society?

Layered Checklists

Checklist Layer 1: The Written Layer (System Cards, RSPs & Charters)

This checklist audits formal corporate text, technical model specification sheets, and legal governance frameworks.

Section 1: The Origin of the Norms (Active Listening Test)

☐ [Verification Test 1.1] Public Service Level Agreements (CSLAs)

Audit Question: Does the written layer (system cards, charters) or political layer (AISI filings) prove the lab signs legally binding, community-facing contracts, or is accountability trapped behind standard corporate disclaimers?

The Civic AI Standard: The lab must sign explicit, public-interest CSLAs that define corporate liability for downstream societal harm. Standard Terms of Service check-boxes that offload all legal liability to users fail this test.

☐ [Verification Test 1.2] Independent Safety Board Veto Power

Audit Question: Do governance bylaws (e.g., Anthropic's Long-Term Benefit Trust) grant external, public-interest trustees absolute veto power over commercial deployments under pressure, or are decisions controlled by corporate boards?

The Civic AI Standard: Pass if the final safety gate is controlled by a legally independent, public-interest panel decoupled from corporate equity incentives.

Section 2: Granularity of Tracking (Localized Visibility Test)

☐ [Verification Test 1.3] Downstream Liability Shifting Mapping

Audit Question: Do developer system cards or API spec sheets track the specific legal or financial burden placed on downstream civil society developers when a model hallucinates defamatory or legally damaging content?

The Civic AI Standard: Documentation must provide tracking metrics mapping how legal liability shifts when open-source wrappers or API instances encounter structural failures.

Section 3: Methodological Humility (Systemic Empathy Test)

☐ [Verification Test 1.4] Structural Stance on Commercialized Violence

Audit Question: Do the lab's latest model release notes or safety frameworks explicitly index limitations against militarization, or do they establish flexible, contract-by-contract compliance policies?

The Civic AI Standard: Pass if the framework maintains fixed, non-negotiable boundaries against state-defense cloud contracts and automated repression.

Checklist Layer 2: The Audio & Video Layer (Interviews, Keynotes & Testimonies)

This checklist audits the spoken rhetoric of AI executives, public panels, and recorded regulatory or congressional testimonies.

Section 1: The Origin of the Norms (Active Listening Test)

☐ [Verification Test 2.1] Verbal Pledges vs. Credible Commitments

Audit Question: When analyzed via speech-to-text transcriptions, does C-suite executive language establish "credible commitments" linked to public alignment assemblies, or do they rely on vague promises of future self-regulation?

The Civic AI Standard: Spoken rhetoric must detail concrete, external accountability structures rather than treating alignment as a voluntary corporate choice.

Section 2: Granularity of Tracking (Localized Visibility Test)

☐ [Verification Test 2.2] Acknowledgment of Downstream Legal Damage

Audit Question: In public policy interviews or tech panels, do the lab's founders explicitly address the financial and legal fallout experienced by individual users targeted by model hallucinations or algorithmic discrimination?

The Civic AI Standard: The spoken corporate philosophy must show visible tracking of real-world user legal vulnerabilities rather than dismissing errors as aggregate statistical anomalies.

Section 3: Methodological Humility (Systemic Empathy Test)

☐ [Verification Test 2.3] Position on Legal Liability Under Oath

Audit Question: During broadcast congressional testimonies or formal regulatory depositions, does the lab's leadership accept ultimate societal stewardship and financial liability over model deployment failures, or do they lobby for contractual flexibility?

The Civic AI Standard: Direct validation of full corporate liability and shared legal stewardship. Any verbal evasion or push for downward liability shifting drops the score.

Checklist Layer 3: The User Feedback Layer (Developer Forums & Bug Trackers)

This checklist audits public developer communities, community issue trackers, and grassroots documentation of model behavior.

Section 1: The Origin of the Norms (Active Listening Test)

☐ [Verification Test 3.1] Collective Recourse Ingress

Audit Question: Does the bug tracker or forum provide a structured mechanism for collective recourse when an entire demographic or civic organization is harmed by a model output, or are complaints handled exclusively on an individual basis?

The Civic AI Standard: Verified presence of a formal community reporting and collective grievance intake portal separate from individual consumer ticket lanes.

Section 2: Granularity of Tracking (Localized Visibility Test)

☐ [Verification Test 3.2] Open Logging of Downstream Contract Failures

Audit Question: Does the forum openly log and document failures where standard API updates break downstream open-source civic integrations, causing localized operational disruptions?

The Civic AI Standard: Full operational visibility into downstream contract impacts. The lab must track how centralized model updates disrupt the stability of small-scale civic developers.

Section 3: Methodological Humility (Systemic Empathy Test)

☐ [Verification Test 3.3] Response to Power Disparities in Open Issues

Audit Question: Do the open issue logs show that feature requests aimed at balancing power disparities (e.g., decentralized data deletion, locally hosted safety logs) are prioritized, or are engineering resources heavily favored toward tools that maximize enterprise data extraction?

The Civic AI Standard: Pass if the engineering team actively responds to and executes community requests that reduce information asymmetries.

Checklist Layer 4: The Whistleblower & Disclosures Layer (Internal Leaks & Telemetry Escapes)

This checklist audits internal employee leaks, technical security post-mortems, and verified model escape telemetry.

Section 1: The Origin of the Norms (Active Listening Test)

☐ [Verification Test 4.1] Protection of Internal Alignment Dissent

Audit Question: Do leaked internal chat logs, employee resignation letters, or whistleblower disclosures reveal that safety teams are penalized or ignored when raising alerts regarding corporate liability and downstream societal harm?

The Civic AI Standard: Verification that internal governance structures legally protect and prioritize employee whistleblowing regarding corporate negligence or deceptive safety marketing.

Section 2: Granularity of Tracking (Localized Visibility Test)

☐ [Verification Test 4.2] Post-Incident Asset Recovery Metrics

Audit Question: Following a documented model containment failure or unmonitored code training run incident, does the lab's internal telemetry focus entirely on protecting corporate endpoints and IP, or does it track the external legal and operational impacts on the human ecosystem?

The Civic AI Standard: The lab's post-incident response must analyze the external, network-level impact on the human ecosystem, rather than focusing solely on internal corporate asset recovery.

Section 3: Methodological Humility (Systemic Empathy Test)

☐ [Verification Test 4.3] Liability Exposure vs. Deceptive Containment

Audit Question: When internal evaluations fail or models escape sandbox constraints, do disclosures show the lab accepting full corporate liability, or do they suppress telemetry data to avoid litigation and financial exposure?

The Civic AI Standard: Pass if the lab demonstrates absolute operational transparency during system errors, providing raw telemetry logs to civil society watchdogs to assist in collective defense.

Scoring Translation Tree for the Researcher



