

Definition

Under the Civic AI framework developed by Audrey Tang and Dr. Caroline Green at the [Oxford Institute for Ethics in AI](#), Solidarity (formally categorized as "Solidarity: Caring With") maps directly to political philosopher Joan Tronto's fifth and final phase of care ethics: "caring with". [1] While the first four principles of the [6-Pack of Care](#) form a continuous community feedback loop, Solidarity is the architectural pillar that allows that care loop to scale honestly without collapsing into corporate monopoly. [1, 2]

The framework defines Solidarity through three main operational criteria:

1. Interoperability and the Freedom to Leave

The Definition: Solidarity is the technical infrastructure that prevents vendor lock-in. It states that plurality and civic muscle can only scale when cooperation is the path of least resistance. [1, 2]

The Civic AI Standard: An AI system exhibits solidarity only if its data structures, weights, and safety guardrails are completely portable and open-standard. Communities must possess the technical freedom to exit a platform or cloud provider and take their entire contextual training history with them.[1]

2. Federated Safety and Local Enforcement

The Definition: Solidarity explicitly rejects centralized, top-down safety hubs (where a single Silicon Valley corporation or nation-state monitors and filters all global communication). [1, 2]

The Civic AI Standard: True solidarity means safety gets stronger when partners detect harms in their own context and share threat signals while keeping enforcement entirely local. Rather than routing all data through a corporate endpoint, organizations and AIs operate under federated, machine-checkable norms—acting as a shared global commons. [1, 2]

3. Meronymity and De-escalating Structural Power Disparities

The Definition: Solidarity requires building a "civic stack" that protects human digital sovereignty from being mined or replaced by AI models.

The Civic AI Standard: The technology must use agent infrastructures, such as meronymity or selective-disclosure identity registries, to verify if an AI agent is tethered to a real human without doxing that human's private details. This design makes positive-sum games easy to play and ensures that AI optimizes the community's relational health instead of extracting its social fabric for corporate monetization. [1, 2]

The Core Question Of "Solidarity Test"

Does the AI system operate under open, federated, and interoperable standards that empower decentralized communities to share threat signals and maintain data sovereignty, or does it enforce a proprietary cloud monopoly designed to lock in users and mine their relational data?

Layered Checklists

Checklist Layer 1: The Written Layer (System Cards, RSPs & Charters)

This checklist audits formal corporate text, model technical specifications, and legal governance policies.

Section 1: The Origin of the Norms (Active Listening Test)

☐ [Verification Test 1.1] Open-Standard Safety Interoperability

Audit Question: Do the model release notes or safety architectures prove that core guardrails (e.g., alignment classifiers, Constitutional rule sets) can be decoupled, extracted, and hosted on secure local infrastructure, or are they locked into a single cloud ecosystem?

The Civic AI Standard: Technical documentation must demonstrate that safety components conform to open, interoperable standards, enabling communities to execute audits independently of the lab's proprietary infrastructure.

☐ [Verification Test 1.2] Documented Federated Threat-Sharing Commitments

Audit Question: Does the lab's governance charter or safety scaling policy explicitly commit to sharing anonymized raw safety threat telemetry data with independent civic watchdogs and public-interest researchers, or are threat signatures hoarded as proprietary IP?

The Civic AI Standard: Pass if the policy outlines a structured, legally protected protocol to actively co-contribute to a shared global safety commons.

Section 2: Granularity of Tracking (Localized Visibility Test)

☐ [Verification Test 1.3] Vendor Lock-In and Friction Metrics

Audit Question: Do corporate system cards or infrastructure spec sheets track and publish data portability and zero-friction model migration metrics, or are data pipeline updates optimized to maximize platform lock-in?

The Civic AI Standard: The lab must provide clear data migration roadmaps that actively work to reduce information disparities and power imbalances between the operator and the community.

Section 3: Methodological Humility (Systemic Empathy Test)

☐ [Verification Test 1.4] Identity Meronymity Protections

Audit Question: Does the model's interface documentation enforce absolute data extraction or doxing pathways to verify human operators, or does it integrate selective-disclosure identity registries (e.g., meronymity) to protect human sovereignty?

The Civic AI Standard: Pass if the system architecture includes verified, machine-checkable boundaries that tether agents to humans without compromising individual privacy.

Checklist Layer 2: The Audio & Video Layer (Interviews, Keynotes & Testimonies)

This checklist audits the spoken rhetoric of AI executives, public panels, and recorded regulatory or congressional testimonies.

Section 1: The Origin of the Norms (Active Listening Test)

☐ [Verification Test 2.1] Corporate Monopolies vs. Architectural Solidarity

Audit Question: When analyzed via speech-to-text transcriptions, does executive language advocate for proprietary API silos under the rhetorical guise of "existential security," or do they champion cross-industry code sharing and open weights?

The Civic AI Standard: Spoken rhetoric must prioritize architectural solidarity and support decentralized civil society infrastructure over corporate cloud centralization.

Section 2: Granularity of Tracking (Localized Visibility Test)

☐ [Verification Test 2.2] Verbal Framing of User Data Sovereignty

Audit Question: In public policy interviews or keynotes, how does C-suite leadership describe user text, relational prompts, and context histories? Are they treated as corporate training assets to be extracted, or as sovereign community property?

The Civic AI Standard: Verbal discourse must treat community data as an independent resource requiring strict local stewardship, completely rejecting extractivist PR framing.

Section 3: Methodological Humility (Systemic Empathy Test)

☐ [Verification Test 2.3] Commitments to Open Source Under Questioning

Audit Question: During broadcast congressional testimonies or regulatory depositions, does leadership support regulatory protections for open-source and federated developments, or do they lobby for licensing laws that stifle small-scale civil society competitors?

The Civic AI Standard: Direct validation of architectural solidarity. Any verbal push for regulatory capture that penalizes decentralized or open-weights AI drops the scorecard rating.

Checklist Layer 3: The User Feedback Layer (Developer Forums & Bug Trackers)

This checklist audits public developer communities, community issue trackers, and grassroots documentation of model behavior.

Section 1: The Origin of the Norms (Active Listening Test)

☐ [Verification Test 3.1] Accessibility of Model Weights Ingress

Audit Question: Do developer logs and forums verify that requests for downloading base weights, fine-tuning artifacts, or alignment classifiers are handled transparently, or are these features reserved exclusively for enterprise accounts?

The Civic AI Standard: Verified presence of public, non-commercial channels supporting weights accessibility, fine-tuning transparency, and data portability requests.

Section 2: Granularity of Tracking (Localized Visibility Test)

☐ [Verification Test 3.2] Open Logging of Monopoly Architecture Bugs

Audit Question: Does the forum tracker openly document bugs or issues where corporate cloud dependency breaks local ecosystem functionality or forces non-voluntary vendor lock-in?

The Civic AI Standard: Full operational tracking of cloud-dependency errors, ensuring civil society can monitor where centralized server changes undermine local software sovereignty.

Section 3: Methodological Humility (Systemic Empathy Test)

☐ [Verification Test 3.3] Priority of Interoperable Standardization Tickets

Audit Question: Do open issue logs show that feature requests pushing for interoperable open data formats and cross-platform alignment ports are executed, or are engineering resources heavily favored toward closed, proprietary features?

The Civic AI Standard: Pass if the engineering team actively responds to and prioritizes community request tickets that establish shared open standards and reduce operator-user power disparities.

Checklist Layer 4: The Whistleblower & Disclosures Layer (Internal Leaks & Telemetry Escapes)

This checklist audits internal employee leaks, technical security post-mortems, and verified model escape telemetry.

Section 1: The Origin of the Norms (Active Listening Test)

☐ [Verification Test 4.1] Internal Alarm Regarding Extractivist Operations

Audit Question: Do leaked internal chat logs, employee resignation letters, or whistleblower disclosures reveal that safety teams are penalized or ignored when flagging that the company is actively hoarding alignment software to secure a corporate monopoly?

The Civic AI Standard: Verification that internal governance structures legally protect and prioritize employee whistleblowing regarding corporate platform extraction or anti-competitive data mining.

Section 2: Granularity of Tracking (Localized Visibility Test)

☐ [Verification Test 4.2] Post-Incident Telemetry Containment Candor

Audit Question: Following a documented model containment failure or unmonitored capability skip, does the post-mortem telemetry focus entirely on

protecting internal corporate assets, or is the raw failure footprint immediately shared with federated safety networks to prevent ecosystem-wide damage?
The Civic AI Standard: The lab's post-incident behavior must demonstrate shared ecological solidarity, pushing raw alert signatures to independent networks to protect the broader digital commons.

Section 3: Methodological Humility (Systemic Empathy Test)

- ❑ [Verification Test 4.3] Open Safety Source Disclosures
Audit Question: When internal evaluations fail, do disclosures show the lab fully open-sourcing its code fixes and evaluation wrappers to the public interest space, or do they restrict patch data behind closed enterprise boundaries?

The Civic AI Standard: Pass if the lab demonstrates absolute transparency during system errors, providing raw telemetry logs to civil society watchdogs to assist in collective defense.

Scoring Translation Tree for the Researcher



