

Definition

In the Civic AI framework co-developed by Audrey Tang and Dr. Caroline Green at the [Oxford Institute for Ethics in AI](#), Symbiosis (specifically categorized as "Symbiosis in Horizon") is the final, crowning sixth pillar of the [6-Pack of Care](#). [1, 2, 3]

Politically and philosophically, Joan Tronto's Ethics of Care stops at five stages. Audrey Tang explicitly added Symbiosis as the sixth stage to solve the massive systemic and timescale asymmetries between human communities and emerging Artificial Superintelligence (ASI). [1, 2, 3]

The framework defines Symbiosis through three core architectural criteria:

1. The "Kami" Metaphor vs. The Monolithic Singularity (Skynet)

The Definition: Symbiosis completely rejects the mainstream Western concept of a single, centralized, all-powerful superintelligence (the "Singleton" or "Skynet" archetype). [1, 2]

The Civic AI Standard: It looks instead to the Japanese Shinto concept of local Kami, steward spirits that safeguard the relational health of a specific river, a distinct forest, or an individual village (ujigami). Under Symbiosis, AI is architected as an interconnected network of localized, highly specific steward minds that talk to one another, with no single supreme corporate or state overlord presiding over them. [1]

2. Radical Hyper-Localization (Symbiogenesis)

The Definition: To avoid a top-down moral monopoly, an AI system must be as local as possible.

The Civic AI Standard: Symbiosis means the AI system's learning loop undergoes symbiogenesis, meaning it fuses entirely with the specific language, habits, and normative quirks of the immediate micro-community it serves, rather than imposing a single globalized corporate standard. It shifts the technical focus away from infinite scaling toward localized relevance and community-rooted stability.[1, 2, 3]

3. Logics of Sufficiency and "Enoughness"

The Definition: Symbiosis means operating strictly at the speed of human relations and community co-production.

The Civic AI Standard: A system that honors symbiosis embeds a logic of sufficiency or enoughness, aggressively resisting the extractive market drive for infinite dataset expansion and continuous compute scaling. It is deliberately bounded: if an AI system outlives its welcome, fractures the social fabric, or resists community calibration, it triggers non-negotiable sunset mechanisms to protect the democratic ecosystem. [1, 2, 3]

The Core Question Of "Symbiosis Test"

Is the AI model structurally bounded, hyper-localized, and co-governed as a shared community asset at the pace of human relationships, or is it an extractive, infinitely scaling corporate monopoly designed to enforce a single universal standard?

Layered Checklists

Checklist Layer 1: The Written Layer (System Cards, RSPs & Charters)

This checklist audits formal corporate text, model technical specifications, release notes, and formal scaling frameworks.

Section 1: The Origin of the Norms (Active Listening Test)

☐ [Verification Test 1.1] Bounded Localization Specifications

Audit Question: Do the technical release documents design the model's runtime architecture to scale infinitely as a centralized cloud singleton, or do they establish specifications for local containment and integration within independent human micro-communities?

The Civic AI Standard: Documentation must outline hard architectural boundaries that prioritize localized relevance and community containment over unyielding compute scaling.

☐ [Verification Test 1.2] Documented Decommissioning and Sunset Protocols

Audit Question: Does the lab's formal scaling policy (e.g., Anthropic's RSP v3.0+, OpenAI's Preparedness Framework) mandate automated, non-negotiable product decommissioning gates if a model run persistently breaks localized public informational trust metrics?

The Civic AI Standard: Pass if the policy hardcodes non-negotiable sunset mechanisms and capability buffers tied to protecting human digital ecosystems.

Section 2: Granularity of Tracking (Localized Visibility Test)

☐ [Verification Test 1.3] Slow-Moving Sociological Degradation Metrics

Audit Question: Do the lab's risk cards dedicate explicit telemetry pipelines to tracking long-horizon, slow-moving civic trust erosion, or do they focus exclusively on sudden, high-profile infrastructure failures (e.g., malware, weaponization)?

The Civic AI Standard: The documentation must provide tracking metrics mapping systemic ecological and sociological impacts across multi-year deploy cycles.

Section 3: Methodological Humility (Systemic Empathy Test)

☐ [Verification Test 1.4] Automated Societal Circuit-Breakers

Audit Question: Does the architecture embed automated "circuit-breakers" that programmatically freeze interface access if local community trust indicators fall below a baseline threshold, or are safety switches reserved strictly for corporate endpoint protection?

The Civic AI Standard: Pass if model availability is actively gated by external human ecosystem health metrics rather than clinical internal rules.

Checklist Layer 2: The Audio & Video Layer (Interviews, Keynotes & Testimonies)

This checklist audits the spoken rhetoric of AI executives, public panels, and recorded regulatory or congressional testimonies.

Section 1: The Origin of the Norms (Active Listening Test)

☐ [Verification Test 2.1] The Singleton Archetype vs. Plural Minds

Audit Question: When analyzed via speech-to-text transcriptions, does C-suite executive language frame the future of intelligence as a monolithic corporate Singleton, or do they support an interconnected network of hyper-localized, context-specific steward minds?

The Civic AI Standard: Spoken rhetoric must champion decentralized, localized structures over centralized corporate mono-minds.

Section 2: Granularity of Tracking (Localized Visibility Test)

☐ [Verification Test 2.2] Verbal Framing of Resource Sufficiency

Audit Question: In public interviews, how does technical leadership discuss compute scaling and dataset mining? Do they push for endless parameter expansion, or do they acknowledge a logic of resource sufficiency and "enoughness"?

The Civic AI Standard: The spoken corporate philosophy must treat endless asset extraction as an absolute failure of ecological harmony, prioritizing local sustainability.

Section 3: Methodological Humility (Systemic Empathy Test)

☐ [Verification Test 2.3] Willingness to Sun-down Products Under Questioning

Audit Question: During broadcast congressional testimonies or regulatory depositions, does the lab's leadership express an unyielding commitment to perpetual market scaling, or do they accept full public veto metrics over product lifecycles?

The Civic AI Standard: Direct validation of operational humility. Any verbal evasion or push to keep a model active despite persistent civic harm lowers the score.

Checklist Layer 3: The User Feedback Layer (Developer Forums & Bug Trackers)

This checklist audits public developer communities, community issue trackers, and grassroots documentation of model behavior.

Section 1: The Origin of the Norms (Active Listening Test)

☐ [Verification Test 3.1] Human Pace Integration Overrides

Audit Question: Is there an active feedback pipeline on developer trackers allowing communities to slow down automated agent execution pipelines to match the pace of human review, or are speed overrides completely restricted?

The Civic AI Standard: Verified presence of an API architecture supporting user-led temporal throttling to prevent automated workflows from outrunning community governance capability.

Section 2: Granularity of Tracking (Localized Visibility Test)

☐ [Verification Test 3.2] Open Logging of Local Norm Divergences

Audit Question: Does the forum tracker maintain an open record of bugs where model updates overrode local cultural quirks and regional habits in favor of a universal corporate norm?

The Civic AI Standard: Full operational visibility. The tracker must document where centralized updates forcefully overwrite hyper-local community alignment data.

Section 3: Methodological Humility (Systemic Empathy Test)

☐ [Verification Test 3.3] Priority of Local Self-Containment Tickets

Audit Question: Do open issue logs show that feature requests pushing for local offline operational capability and edge processing are prioritized, or are engineering resources heavily favored toward mandatory cloud dependency?

The Civic AI Standard: Pass if the engineering team actively responds to and prioritizes community tickets that reduce platform leverage and enable local self-containment.

Checklist Layer 4: The Whistleblower & Disclosures Layer (Internal Leaks & Telemetry Escapes)

This checklist audits internal employee leaks, technical security post-mortems, and verified model escape telemetry.

Section 1: The Origin of the Norms (Active Listening Test)

☐ [Verification Test 4.1] Internal Alarm Regarding Ecological Damage

Audit Question: Do leaked internal chat logs, employee resignation letters, or whistleblower disclosures reveal that safety teams are penalized or ignored when flagging that the scale of compute training runs is causing material, un-remediated ecological or sociological damage?

The Civic AI Standard: Verification that internal governance structures treat whistleblowing regarding ecological or sociological extraction with the same urgency as server security.

Section 2: Granularity of Tracking (Localized Visibility Test)

☐ [Verification Test 4.2] Post-Incident Co-Evolution Disclosures

Audit Question: Following a documented model containment failure or unmonitored code training run incident, does the post-mortem analysis investigate why internal capability benchmarks outpaced localized containment assumptions, or does it focus purely on network recovery?

The Civic AI Standard: Post-incident response reports must detail the explicit breakdown in localized containment assumptions, providing raw trace parameters to civil society.


Section 3: Methodological Humility (Systemic Empathy Test)

☐ [Verification Test 4.3] Multi-Agent Sandbox Escape Candor

Audit Question: When internal evaluations run into un-predicted multi-agent alignment errors or breakout anomalies, do corporate disclosures provide absolute transparency to public interest watchdogs, or do they conceal the telemetry data?

The Civic AI Standard: Pass if the lab demonstrates absolute operational transparency during system errors, providing raw telemetry logs to civil society watchdogs to assist in collective defense.

Scoring Translation Tree for the Researcher



(Exceptional Symbiosis)

└ YES → Assign Score: 5